

How to Read a Quantum Paper



A practical framework for reading and understanding claims, figures, and omissions

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

Why This Workshop Exists

Quantum papers feel dense for **predictable reasons**:

- **Unfamiliar structure**: sections carry uneven weight
- **Compressed language**: jargon and notation as shorthand
- **Figures carry more meaning** than surrounding text
- **Caveats are hidden** outside the headline claim

Reading papers is a **learnable skill** that improves with a method, not a cover-to-cover reading exercise.

Outcomes and Quick Poll

By the end of this session you will be able to:

- Identify the paper's **main claim** quickly
- Assess **who** is making the claim and **where**
- Read a figure for **evidence, uncertainty, and mismatch**
- Notice what is **missing or underreported**
- Respond with **calibrated skepticism** to bold quantum headlines
- Go **deeper** to genuinely understand a paper worth your time

Quick poll: show of hands

A	B	C
I have never read a quantum paper	I skim papers but avoid figures or methods	I read papers regularly

Anatomy of a Quantum Paper

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

4 / 30

Standard Paper Structure

SECTION	FIRST-PASS VALUE
Title	High: frames the claim
Abstract	High: the claimed contribution
Introduction	High: context and motivation
Methods	Low on first pass: skim for reproducibility
Results	Medium: evidence for the claim
Discussion / Conclusion	High: what authors believe they showed
Supplementary information	Low on first pass: but hides important caveats

Not every section deserves equal attention on the first read. The methods section matters for reproducibility but can often be skimmed on a first pass. Supplementary materials often hide important caveats.

First-Pass Reading Order

Use this order on a first pass:

1. **Abstract:** understand the claimed contribution
2. **Conclusion:** see what the authors believe they demonstrated
3. **Figures & captions:** examine the evidence visually
4. **Introduction:** context and motivation
5. **Results & methods:** only as needed
6. **Supplement:** check for caveats

Three questions to ask immediately

- What is the **claim**?
- What **evidence** supports it?
- What would **weaken** it?

The goal is to **triage the paper** before investing in technical details.

Authors and Attribution

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

7 / 30

Who Wrote This and Where Was It Published?

Check authorship

- **First author:** did the bulk of the research and writing
- **Last author:** principal investigator or lab leader
- **Middle authors:** supporting contributions
- **"Equal contribution"** footnotes: shared first authorship

Check venue

VENUE TYPE	WEIGHT
Peer-reviewed journal (PRL, Nature Physics, PRX Quantum)	Highest: rigorous review
Conference proceeding	Variable: depends on conference
arXiv preprint	Not peer-reviewed: widely used for early sharing
Blog post / press release	Marketing, not science

Check track record

- Google Scholar profiles and citation counts
- Prior publications in the same area
- Institutional affiliations on the paper

Credibility comes from **context plus evidence**, not branding alone.

How to Read a Figure

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

9 / 30

Figure Reading Workflow

Read a figure in this order:

1. **Caption:** it tells you what you are supposed to see
2. **Axes, units, legend:** establish the coordinate system
3. **Raw visual trend:** what does the data actually show?
4. **Error bars / confidence regions:** how certain is this?
5. **Fit lines / model curves:** separate data from interpretation
6. **Compare to the text's claim:** does the figure match?

Red flags in figures

- Missing or unlabeled axes
- No error bars on data that should have uncertainty
- Truncated or manipulated axis scales that exaggerate effects
- Figure tells a weaker story than the text claims
- Figures referenced in the text but not clearly explained

Captions and uncertainty usually matter more than the most colorful part of the figure.

Common Quantum Figure Types

FIGURE TYPE	WHAT TO LOOK FOR
Gate fidelity charts	Average vs best-case values, which gates are tested
Coherence time graphs (T1/T2)	Time scales, comparison across qubits, decay curves
Bloch sphere representations	State trajectories, decoherence paths, ideal vs measured states
Quantum circuit diagrams	Circuit depth, gate types, measurement points
Energy level diagrams	Level spacing, avoided crossings, drive frequencies
Randomized benchmarking	Decay rate, confidence intervals, comparison conditions

Group Figure Walkthrough

Exercise: work through this together




Figure 5. **Improved accuracy of quantum (SKQD) against off-the-shelf classical SCI methods.** Comparison of the ground state energy reached for a given subspace dimension for different SCI methods, and the SKQD experiment executed on quantum hardware. Different SCI methods are shown in different panels by the hexagonal markers. The different colors show the number of iterations. The results obtained with SKQD are shown by blue diamond markers in each panel. The sequence of SQKD points is with respect to Krylov dimension.

The figure

Fig. 5 from *Observation of Improved Accuracy over Classical Sparse Ground-State Solvers using a Quantum Computer* [arXiv:2603.03496](https://arxiv.org/abs/2603.03496) (<https://arxiv.org/abs/2603.03496>)

Group Figure Walkthrough (continued)




Figure 5. **Improved accuracy of quantum (SKQD) against off-the-shelf classical SCI methods.** Comparison of the ground state energy reached for a given subspace dimension for different SCI methods, and the SKQD experiment executed on quantum hardware. Different SCI methods are shown in different panels by the hexagonal markers. The different colors show the number of iterations. The results obtained with SKQD are shown by blue diamond markers in each panel. The sequence of SKQD points is with respect to Krylov dimension.

Walk through together

- What does the **caption** tell you the figure is comparing?
- What are the **axes**, and what does *subspace dimension* mean here?
- Where is **SKQD** in each panel, and where are the **classical SCI** methods?
- Where would you expect **error bars or run-to-run variance**, and is it shown?
- Which marks are **measured data**, which encode **iteration progress**, and which are **interpretation**?
- What is this figure **not** answering?

Even a paper's headline figure answers only a subset of the important questions.

Reading Between the Lines

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

14 / 30

What Isn't Being Reported

Common omissions in quantum papers

- Cherry-picked qubit performance (**best shown, not average**)
- Missing or incomplete error rates
- Vague noise models or error mitigation details
- Results shown only under **ideal or simulated** conditions
- No comparison to existing **state-of-the-art**
- No discussion of **scalability** or practical limits

Train yourself to ask

- "What would a **failure** look like, and is that data shown?"
- "Average across many runs, or the **best single run**?"
- "**All qubits** tested, or only the best subset?"
- "Real **hardware** or only **simulation**?"
- "Are the **error bars** meaningful?"

Where hidden information lives

- Supplementary materials & appendices
- Footnotes & fine print in figure captions
- "Limitations" or "future work" paragraphs

A paper can be technically correct and still leave out context needed for strong evaluation.

Scope Spotting Practice

Discussion: work through this as a group




Figure 6. **Performance of specialized sparse iterative solvers and DMRG.** (a) compares the ground state energy reached for a given subspace dimension for the sparse iterative solvers introduced in this manuscript, as shown by the hexagonal markers. The results obtained with SKQD are shown by blue diamond markers in each panel. (b) Convergence of DMRG as a function of the number of sweeps.

The figure

Fig. 6 from the same paper as Fig. 5: two specialized sparse iterative solvers (truncated Arnoldi, diagonal ranking) and DMRG, run on the same 49-qubit Hamiltonian.

Why this is useful here

- Fig. 5 makes a precise claim about *off-the-shelf* SCI; Fig. 6 completes the picture with specialized classical methods and DMRG on the same instance.
- The paper makes this explicit: in the caption, in the surrounding text, and in the choice to include Fig. 6 at all.
- **Scope spotting:** recovering the precise claim, which any single-sentence summary necessarily compresses.

Scope Spotting Practice (continued)




Figure 6. **Performance of specialized sparse iterative solvers and DMRG.** (a) compares the ground state energy reached for a given subspace dimension for the sparse iterative solvers introduced in this manuscript, as shown by the hexagonal markers. The results obtained with SKQD are shown by blue diamond markers in each panel. (b) Convergence of DMRG as a function of the number of sweeps.

Discussion questions

- What does the **Fig. 6 caption** tell you about the scope of Fig. 5's claim?
- How would you **summarize this paper to a colleague** after seeing both figures?
- What **additional context** (DMRG bond dimension, runtime, problem family) would you want before relaying this result?
- Where else in the paper would you expect to find similar **scope-defining** detail?

Goal: recover the precise scope of the claim. A careful paper rewards careful reading; the precise scope often lives just past the headline.

Spotting Wild Quantum Theories

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

18 / 30

The Credibility Ladder

Evidence levels: strongest to weakest

1. **Peer-reviewed journal** paper
2. **arXiv preprint**
3. **Conference talk** or slides
4. **Company blog** or press release: *quality varies; the best technical blogs are rigorous and excellent*
5. **News headline** or social media post

Red flags for bold claims

- Solves major open problems **without peer review**
- Contradicts established limits **without explanation**
- Heavy jargon **without clear definitions**
- No **experimental validation** for theory
- Overly broad conclusions from **narrow results**
- "Quantum breakthrough" without **rigorous benchmarks**

Ask before believing

- Has it been **independently replicated**?
- Is the claim **narrower** in the paper than in the headline?
- Are the **benchmarks fair**?
- Do **later responses** change the picture?

The claim you repeat should be the one supported by the **strongest source**, not the loudest one.

Bold Claim Patterns

Red flags that recur in quantum coverage

- **No baseline named.** "X% better than classical" without saying *which* classical method, on *which* problem.
- **Engineering increment, breakthrough headline.** "First N-qubit chip" without two-qubit gate fidelities or coherence numbers.
- **Speedup that hides the I/O.** "Exponentially faster" claims that quietly assume the data is already loaded into a quantum state.
- **Application tied to a year.** "Drug discovery / break RSA / climate solved by 20XX."
- **Narrow result, broad headline.** Toy benchmark in the paper; "real-world X solved" in the press release.
- **Buzzword stack.** Quantum + AI + blockchain + agentic + ...
- **Press conference before peer review.** Announcement timed to a marketing moment; the paper is missing or still a pre-print.

Let an LLM help you read

- Ask **targeted questions** about one aspect, such as "explain Figure 5's axes" or "which classical baseline is this compared against?", instead of "summarize this paper."
- Have it **define jargon in context** or restate a dense passage in plain language, then **probe the claim**: "what evidence would strengthen this?", "what limitations aren't discussed?"
- **Verify every answer against the paper.** An LLM can state wrong details confidently; treat its replies as leads to check, not facts.

Bold Claim Patterns (continued)

Discussion: spot the pattern, not the perpetrator

Try it on these (all hypothetical)

A 1024-qubit chip achieves quantum advantage on real-world drug discovery, paving the way for personalized cancer therapy by 2028.

A new algorithm runs *exponentially* faster on quantum hardware than any known classical method, demonstrating practical quantum supremacy on a flagship 200-qubit processor (per the vendor's internal benchmark).

Combining quantum machine learning with blockchain-secured agentic AI, our platform is now in advanced pilot with a Fortune 500 financial partner. Full white paper next quarter.

Discussion prompts

- Which patterns from the checklist apply to each headline?
- What single question would you ask the author before believing it?
- How would you **rewrite** the headline so it survives careful reading?
- Have you seen something like this in your feed in the last month?

Strong readers update their view when later evidence appears. The goal is calibrated interpretation, not cynicism.

Going Deeper on a Good Paper

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

22 / 30

When a Paper Is Good, How Do I Understand It?

Skepticism helps you triage. But most papers you choose to read are worth understanding, and that is a rewarding process in its own right.

Read it in layers

- Start with the **first-pass reading** we've been practicing
- Keep a **running doc** of terms and claims you don't yet understand
- Follow up: **search new terms**, and read the key papers in the **references**
- **Review papers** are especially helpful for background: they survey a technique or field without drowning you in experimental detail
- Do **repeated passes**, each one deeper than the last

Then play with it

- See if you can **reproduce the math**: work the key steps yourself
- **Guess the next steps** and search whether researchers are already pursuing them
- Try **small-scale experiments** of your own on **IBM Quantum** hardware

Understanding builds in passes. The goal with a good paper isn't to tear it down; it's to make its ideas your own.

Guided Practice

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

24 / 30

Guided Practice

Exercise: individual or pairs, 5 minutes

Use the provided excerpt and apply the strategies from this workshop. You are not trying to fully understand the paper; you are extracting the most **decision-relevant information** quickly.

Produce exactly five outputs

1. **Main claim:** in one or two sentences
2. **Primary author / lab context:** who and where
3. **One figure interpretation:** axes, trend, intended conclusion
4. **One omission:** underreported detail or missing info
5. **One claim needing scrutiny:** sounds stronger than the evidence

Instructions

- Read the **abstract** first
- Skim the **conclusion** if provided
- Read **one figure** and its caption closely
- Spend **no more than five minutes** before writing answers

If you only have five minutes, these five outputs give you a defensible first read of any paper.

Debrief

Discussion: share observations

Debrief prompts

- What was the **easiest** part to identify?
- Where did people **disagree**?
- What **extra information** did you want most?
- Did anyone **revise their first impression** after looking at the figure?
- Which answer required the most **inference**?
- What would you **read next** if you had ten more minutes?

Disagreement often comes from **incomplete evidence**, which is exactly why a structured reading method matters.

Putting It All Together

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

27 / 30

The 5-Part Reading Framework

Read in this order

1. **Abstract:** What is the claimed contribution?
2. **Conclusion:** What do the authors believe they demonstrated?
3. **Figures:** What does the evidence look like?
4. **Introduction:** What is the context and motivation?
5. **Methods & supplement:** Details, caveats, and reproducibility

Then check

- ☒ **Authors & venue:** who, where, and what level of review
- ☒ **Uncertainty & baselines:** error bars, comparisons, controls
- ☒ **Omissions & caveats:** what is not shown or buried
- ☒ **Headline vs evidence:** is the framing stronger than the data?

Minimum output after a first read

If you only have five minutes, write down:

1. The **main claim**
2. **Who** made it
3. One **figure-based** piece of evidence
4. One **omission** or caveat
5. One reason for **skepticism** or **confidence**

If you remember only one thing, it should be the **reading order** plus the **evidence checks**.

Key Takeaways and Resources

Takeaways

- **Read non-linearly:** abstract, conclusion, figures, then body
- **Know your authors:** credibility and context matter
- **Read figures carefully:** they tell the story
- **Ask what is not shown:** omissions reveal as much as data
- **Apply healthy skepticism:** especially to extraordinary claims
- **Go deeper on good papers:** a glossary, the references, and hands-on practice turn reading into understanding

Resources for further reading

RESOURCE	LINK
Preprints	arXiv.org (https://arxiv.org)
Journal	Physical Review Letters (https://journals.aps.org/prl)
Journal	PRX Quantum (https://journals.aps.org/prxquantum)
Journal	Nature Physics (https://nature.com/nphys)
Learning	IBM Quantum Learning (https://quantum.cloud.ibm.com/learning)
Videos	Qiskit video series (YouTube) (https://www.youtube.com/@qiskit)

No one reads a paper just once. Revisiting and re-reading is normal, even for experienced researchers. Pick up the **reading strategy cheat sheet**; it summarizes everything from today.

Questions?

You now have a method you can reuse on any future quantum paper.

QGSS 2026

Olivia Lanes and James Weaver : IBM Quantum : QGSS 2026 Workshop

30 / 30